

金融大数据处理技术 实验2

171870604 袁满杰

1.环境部署

为了方便之后多机的部署，采用Docker配置Hadoop环境，构建Dockerfile如下：

```
1 FROM ubuntu
2 MAINTAINER while
3 RUN mkdir /usr/local/java \
4     && apt-get update\
5     && apt-get install -y openssh-server vim\
6     && ssh-keygen -t rsa -f ~/.ssh/id_rsa -P ''\
7     && cat /root/.ssh/id_rsa.pub >> /root/.ssh/authorized_keys\
8     && sed -i 's/PermitEmptyPasswords yes/PermitEmptyPasswords
no /' /etc/ssh/sshd_config\
9     && sed -i 's/PermitRootLogin without-
password/PermitRootLogin yes /' /etc/ssh/sshd_config\
10    && echo " StrictHostKeyChecking no" >>
/etc/ssh/ssh_config\
11    && echo " UserKnownHostsFile /dev/null" >>
/etc/ssh/ssh_config\
12    && echo "root:123456" | chpasswd
13
14 ADD jdk-8u221-linux-x64.tar.gz /usr/local/java/
15 ADD hadoop-3.2.1.tar.gz /
16
17 ENV JAVA_HOME /usr/local/java/jdk1.8.0_221
18 ENV JRE_HOME ${JAVA_HOME}/jre
19 ENV HADOOP_HOME /hadoop-3.2.1
20 ENV CLASSPATH .:${JAVA_HOME}/lib:${JRE_HOME}/lib
21 ENV PATH ${JAVA_HOME}/bin:${HADOOP_HOME}/bin:${HADOOP_HOME}/sbin:$PATH
22
23 CMD [ "sh", "-c", "service ssh start; bash"]
```

其中第三行的 `RUN` 安装vim和ssh服务并配置密钥（这里我同上一个实验一样使用配置本机免密访问本机后构建镜像，从而使得该镜像的所有容器有相同的RSA公私钥实现相互免密访问）；

第14~15行 `ADD` 拷贝宿主机上的jdk和Hadoop安装包进行解压/安装；

最后设置环境变量及启动ssh服务。

以此为基本Hadoop环境的镜像构建容器进行以下三个实验。

2. 单机模式运行Hadoop样例程序

首先运行结果如下：

```

2019-10-02 06:19:02,752 INFO mapreduce.Job: Job job_local269050207_0002 running in uber mode : false
2019-10-02 06:19:02,752 INFO mapreduce.Job: map 100% reduce 100%
2019-10-02 06:19:02,753 INFO mapreduce.Job: Job job_local269050207_0002 completed successfully
2019-10-02 06:19:02,758 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=1336130
        FILE: Number of bytes written=3348954
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
    Map-Reduce Framework
        Map input records=1
        Map output records=1
        Map output bytes=17
        Map output materialized bytes=25
        Input split bytes=117
        Combine input records=0
        Combine output records=0
        Reduce input groups=1
        Reduce shuffle bytes=25
        Reduce input records=1
        Reduce output records=1
        Spilled Records=2
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=0
        Total committed heap usage (bytes)=2862612480
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=123
    File Output Format Counters
        Bytes Written=23
root@a103cd11c7cc:/hadoop-3.2.1# cat output/*
1      dfsadmin
root@a103cd11c7cc:/hadoop-3.2.1#

```

其次，具体步骤上.....按照实验手册上教程来就好了，唯一遇到的两个问题：

- Hadoop运行时不认JAVA_HOME环境变量：

```
ERROR: JAVA_HOME is not set and could not be found.
```

- 貌似不允许root运行Hadoop：

```
ERROR: Attempting to operate on hdfs namenode as root
ERROR: but there is no HDFS_NAMENODE_USER defined. Aborting operation.
```

解决方案：

通过在hadoop-env.sh配置文件中加入以下语句可以顺利解决(顺便也为之后启动yarn配置了这个)：

```

1 export JAVA_HOME=/usr/local/java/jdk1.8.0_221
2 export HDFS_NAMENODE_USER="root"
3 export HDFS_DATANODE_USER="root"
4 export HDFS_SECONDARYNAMENODE_USER="root"
5 export YARN_RESOURCEMANAGER_USER="root"
6 export YARN_NODEMANAGER_USER="root"

```

最后顺利得到运行结果

3. 伪分布模式运行Hadoop样例程序

首先运行结果如下：

```

bytes written: 200
root@d2cb5bfla4df:/hadoop-3.2.1# hdfs dfs -get output2 output
2019-10-03 03:42:13,197 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
root@d2cb5bfla4df:/hadoop-3.2.1# cat output/*
5 dfs.audit.logger
3 dfs.logger
3 dfs.server.namenode,
2 dfs.audit.log.maxbackupindex
2 dfs.sh
2 dfs.audit.log.maxfilesize
1 dfsadmin
1 dfs.replication
1 dfs.namenode.servicerpc
1 dfs.namenode.rpc
1 dfs.log
1 dfs.http.address
1 dfs.namenode.ec.policies.max.cellsize
root@d2cb5bfla4df:/hadoop-3.2.1#

```

其次，在创建容器时使用 `-p` 参数指定端口映射从而可以访问Hdfs的ui网页：

```
1 | $ docker run -it -p 8000:50090 hadoop
```

之后也是具体步骤上参照实验手册来配置文件就好了；

其中在修改好配置文件之后我将容器通过 `docker commit` 保存为镜像，这样以后遇到问题想重来就可以直接从这一镜像生成容器进行操作；

其中同样遇到两个问题：

1.宿主机上无法访问Hdfs的ui网页：

经过分析，是由于默认情况下该端口开放在容器中localhost:50070，不对外开放，因此只有容器内部可以访问；

解决方法是在hdfs-site.xml中增加配置：

```

1 | <property>
2 |   <name>dfs.http.address</name>
3 |   <value>0.0.0.0:50070</value>
4 | </property>

```

从而将其改为0.0.0.0:50070，如此便可以在宿主机上、甚至从外网访问宿主机映射后的8000端口进入Hdfs的ui页面了。

2. 重新格式化NameNode之后无法启动DataNode:

遇到之后要提到的第三个问题时，我多次尝试过从头重新再来，因此重新格式化了NameNode，但在没有删除DataNode的情况下重新格式化NameNode会导致其ID与DataNode不同，从而在启动HDFS时会无法启动DataNode（但start-dfs.sh不会报错，在之后拷贝文件进HDFS时才会出错，通过jps命令发现DataNode没有启动）

解决方法有两种：

1. 在HDFS开启状态下进入缓存文件保存地址（默认为/tmp下），然后从name/current/VERSION文件中将第三行的ID复制下来，修改data/current/VERSION中同样位置的ID为一样的，之后就可以正常启动DataNode了
2. 直接删除整个dfs文件夹，让系统重新生成；

3. 显示map0%,reduce 100%之后看似正常结束，但输出结果为空

这个问题卡了我一天.....显示上诡异的map0%,reduce 100%，任务也SUCCESS完成,但是结果为空

```

2019-10-03 03:46:19,751 INFO mapreduce.Job: Job job_local2081552833_0002 running in uber m
ode : false
2019-10-03 03:46:19,752 INFO mapreduce.Job: map 0% reduce 100%
2019-10-03 03:46:19,752 INFO mapreduce.Job: Job job_local2081552833_0002 completed success
fully
2019-10-03 03:46:19,768 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=633251
        FILE: Number of bytes written=1676052
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=0
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=5
        HDFS: Number of bytes read erasure-coded=0
    Map-Reduce Framework
        Combine input records=0
        Combine output records=0
        Reduce input groups=0
        Reduce shuffle bytes=0
        Reduce input records=0
        Reduce output records=0
        Spilled Records=0
        Shuffled Maps=0

```

最后翻去查logs文件夹下的运行日志，发现其中问题所在：

```

2019-10-03 03:33:42,249 INFO org.apache.hadoop.ipc.Server: IPC Server handler 6 on
default port 9000, call Call#67 Retry#0
org.apache.hadoop.hdfs.protocol.ClientProtocol.getBlockLocations from 172.17.0.2:45456:
java.io.FileNotFoundException: Path is not a file: /user/root/input/shellprofile.d

```

原因在于将整个 `etc/hadoop` 目录拷贝到HDFS中时其中有一个 `shellprofile.d` 文件夹（猜测为新版Hadoop中加入的文件夹，因此手册上示例中没问题）也被移入其中，（若单机情况下用cp命令不加参数会提示文件夹无法复制从而跳过该项，之后便可以正常运行），而grep命令时无法对文件夹处理，因此map引发错误中断，reduce对空输出进行处理最终输出也为空（可以看出mapreduce的稳定性很强.....即使map程序出错也能顺利进行reduce,输出Success...）

所以只要将input文件夹中这个子文件夹删除掉就可以正常运行了

4. 集群模式运行Hadoop样例程序

运行结果如下：

```

File Output Format Counters
    Bytes Written=107
root@4f1a3b914bc7:/# hdfs dfs -cat output2/*
2019-10-03 14:33:03,065 INFO sasl.SaslDataTransferClient: SASL encryption trust check: loc
alHostTrusted = false, remoteHostTrusted = false
1      dfsadmin
1      dfs.replication
1      dfs.namenode.secondary.http
1      dfs.namenode.name.dir
1      dfs.datanode.data.dir
root@4f1a3b914bc7:/#

```

在其ui界面中可以看到任务完成情况如下(第一条)：

Cluster Metrics																						
Apps Submitted	3	Apps Pending	0	Apps Running	3	Apps Completed	0	Containers Running	0 B	Memory Used	40 GB	Memory Total	0 B	Memory Reserved	0	VCores Used	40	VCores Total	0	VCores Reserved		
Cluster Nodes Metrics																						
Active Nodes	5	Decommissioning Nodes	0	Decommissioned Nodes	0	Lost Nodes	0	Unhealthy Nodes	0	Rebooted Nodes	0	Shutdown Nodes										
Scheduler Metrics																						
Scheduler Type	Capacity Scheduler				Scheduling Resource Type	(memory-mb (unit=Mi), vcores)				Minimum Allocation	<memory:1024, vCores:1>				Maximum Allocation	<memory:8192, vCores:4>				Maximum Cluster Application Priority	0	
Show	20	▼ entries																		Search		
ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Reserved CPU V-Cores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes		
application_1570106730474_0003	root	grep-sort	MAPREDUCE	default	0	Thu Oct 3 22:32:24 +0800 2019	Thu Oct 3 22:32:24 +0800 2019	Thu Oct 3 22:32:45 +0800 2019	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0	<div></div>	History	0		
application_1570106730474_0002	root	grep-search	MAPREDUCE	default	0	Thu Oct 3 22:31:55 +0800 2019	Thu Oct 3 22:31:56 +0800 2019	Thu Oct 3 22:32:21 +0800 2019	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0	<div></div>	History	0		
application_1570106730474_0001	root	word-count	MAPREDUCE	default	0	Thu Oct 3 20:47:01 +0800 2019	Thu Oct 3 20:47:03 +0800 2019	Thu Oct 3 20:47:27 +0800 2019	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0	<div></div>	History	0		
Showing 1 to 3 of 3 entries																				First Previous 1 Next Last		

首先从之前配置好的Hadoop镜像建立一个临时容器，进入后安装实验手册改写配置文件，这里worker设为h1,h2,h3,h4,h5;主节点设为h0；此外在之前提到的hadoop-env.sh中修改也需加入其中；

之后将修改好配置文件的容器通过 `docker commit` 设为镜像hadoop_3，并编写一下shell脚本据此生成一批容器：

```
1  #! /bin/bash
2
3  for i in {1..5}
4  do
5      docker run -dit --name=h$i --net hadoop-net hadoop_3
6  done
7  docker run -dit --name=h0 -p 8000:50090 -p 8001:8088 -p 8002:19888 --net hadoop-net hadoop_3
```

其中for循环建立了一批worker，加入 `hadoop-net` 虚拟网络中；最后一行将HDFS、yarn的ui端口进行了映射建立主节点

然后通过 `docker attach h0` 进入其bash界面按照实验手册进行后续操作，运行得到结果。

这里只遇到一个问题：

运行时yarn报错：

大概长这样：

```
[2019-04-16 06:25:41.902]Container exited with a non-zero exit code 1. Error file: prelaunch.err.
Last 4096 bytes of prelaunch.err :
Last 4096 bytes of stderr :
Error: Could not find or load main class org.apache.hadoop.mapreduce.v2.app.MRAppMaster
```

Please check whether your etc/hadoop/mapred-site.xml contains the below configuration:

```
<property>
  <name>yarn.app.mapreduce.am.env</name>
  <value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
<property>
  <name>mapreduce.map.env</name>
  <value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
<property>
  <name>mapreduce.reduce.env</name>
  <value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
```

https://blog.csdn.net/qq_36318271

只需按照其解释修改 `mapred-site.xml` 加入以下配置即可：

```
1 <property>
2   <name>yarn.app.mapreduce.am.env</name>
3   <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
4 </property>
5 <property>
6   <name>mapreduce.map.env</name>
7   <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
8 </property>
9 <property>
10  <name>mapreduce.reduce.env</name>
11  <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
12 </property>
```